

Final Data Management Project

Christina Chang

12/8/2017

1. Data Preparation

Load Packages

```
# read data
library(foreign)
library(readr)

# clean & manipulate data
library(tidyverse)
library(dplyr)
library(reshape2)

# plots
library(ggplot2)

# time series
library(tseries)
library(dynlm)
library(urca)

# forecast
library(forecast)
library(scales)

# knit
library(knitr)
library(rvest)
```

Import Data

By visual inspection of the file, we don't want R to read the first two rows. NA argument to specify that any blanks ("") or asterisks ("*") would be considered missing data. Note: the asterisks are specified to be redacted information, based on the UNHCR website.

```
setwd("/Users/Berlin/Desktop/HertieDataScience/final project")

d1 <- read_csv("unhcr_popstats_people of concern.csv", skip = 2, na = c("", "*"))
d2 <- read_csv("unhcr_popstats_refugee status.csv", skip = 2, na = c("", "*"))

View(d1)
View(d2)
```

Tidy Data

The UNHCR data is an unbalanced panel dataset. After tidying the data, we have created a balanced panel dataset with all entities (countries) observed in all years (from 2000 to 2016).

Steps to clean and manipulate data:

```
str(d1)
# Year correctly structured as integer
# Country destination and origin correctly structured as character
# All other variables should be numeric

d1[4:11] <- lapply(d1[4:11], as.numeric)

summary(d1) # years from 1951 to 2016

str(d2)
# Year correctly structured as integer
# Country destination and origin correctly structured as character
# Do not need RSD procedure type information
# All other variables should be numeric

d2[4] <- NULL

d2[4:13] <- lapply(d2[4:13], as.numeric)

summary(d2) # years from 2000 to 2016

df <- merge(d1, d2, by = c("Year", "Country / territory of asylum/residence", "Origin"))

View(df)
str(df)
summary(df) # after merge, data before the year 2000 drops out.

# identify missing data
apply(df, 2, function(x) sum(is.na(x)))
```

2. Exploratory Data Analysis

People of Concern:

```
# subset for only PoC category counts by year

PoC_count <- df[c(1,4:10)]

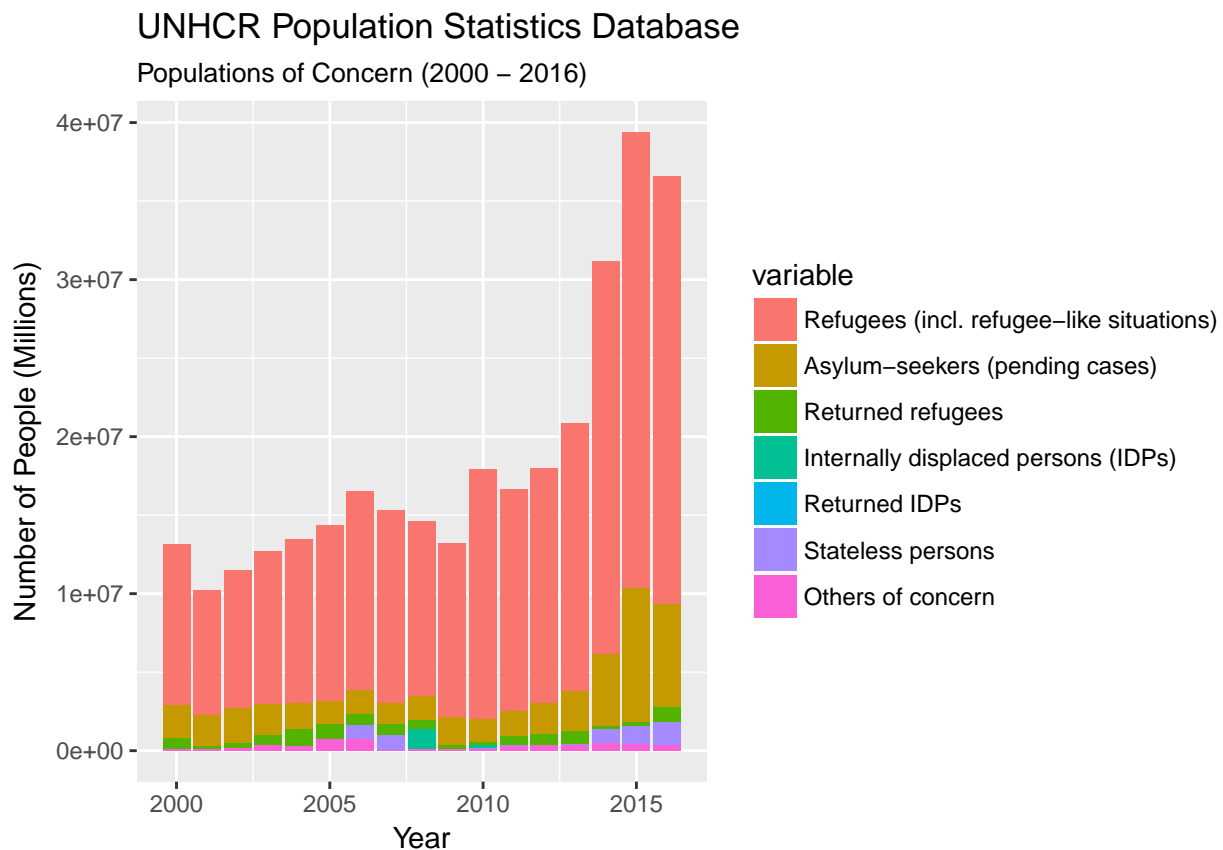
PoC_count <- melt(PoC_count, id=c("Year"))

str(PoC_count)

## 'data.frame': 813456 obs. of 3 variables:
## $ Year : int 2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 ...
## $ variable: Factor w/ 7 levels "Refugees (incl. refugee-like situations)",...: 1 1 1 1 1 1 1 1 1 1 1 .
## $ value : num NA NA 9 507 2 5 NA 1 5 20 ...

plot1 <- ggplot(PoC_count, aes(Year, value, na.rm = TRUE)) +
  geom_bar(aes(fill=variable), stat="identity") +
  labs(title="UNHCR Population Statistics Database",
       subtitle="Populations of Concern (2000 - 2016)",
       x="Year",
       y="Number of People (Millions)")

plot1
```



Percent Change in Total Population by “People of Concern”

```
Year_Pop <- aggregate(df$`Total Population`, by=list(Year = df$Year), FUN=sum, na.rm = TRUE)

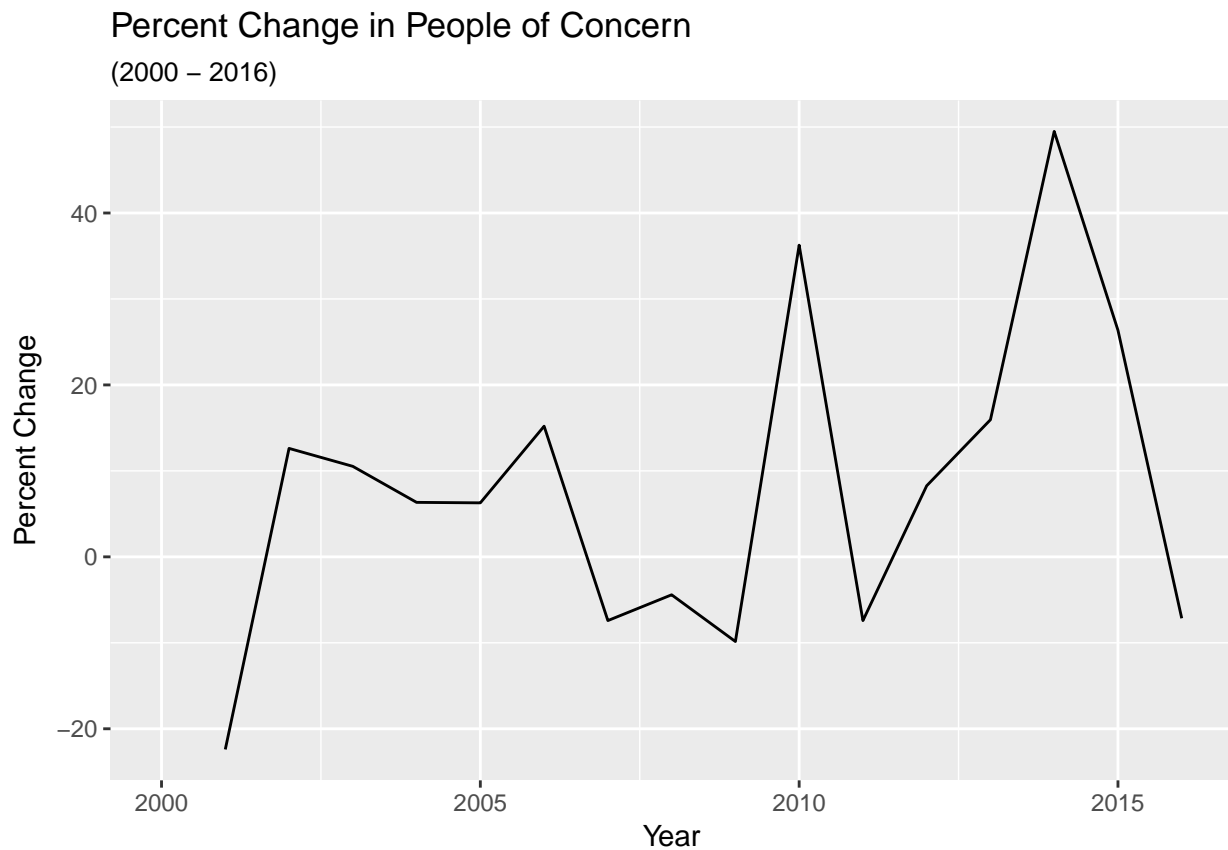
Year_Pop$rate <- NA

Year_Pop$rate[which(Year_Pop$Year>2000)] = 100*(diff(Year_Pop$x)/Year_Pop[-nrow(Year_Pop),]$x)

View(Year_Pop)

plot2 <- ggplot(Year_Pop, aes(x= Year, y= rate)) + geom_line() +
  labs(title="Percent Change in People of Concern",
        subtitle="(2000 - 2016)",
        x="Year",
        y="Percent Change")

plot2
```



Top Countries of Destination

```
destination_country_total <- df %>%
  group_by(`Country / territory of asylum/residence`, Year) %>%
  summarise(Total = sum(`Total Population`))

View(destination_country_total)

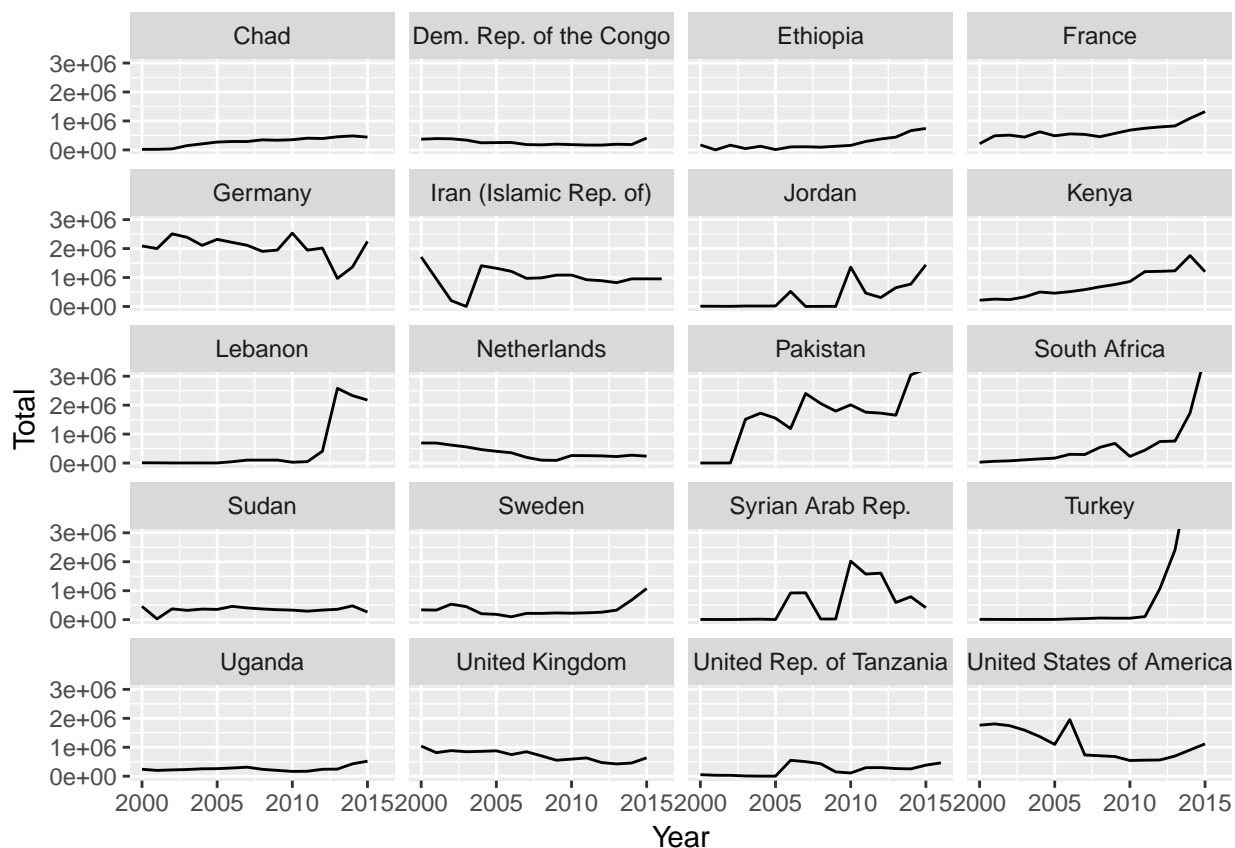
top_destcountries <- destination_country_total %>%
  group_by(`Country / territory of asylum/residence`) %>%
  summarise(Total = sum(Total, na.rm = TRUE)) %>%
  top_n(20)

View(top_destcountries)

top_destcountries2 <- as.character(top_destcountries$`Country / territory of asylum/residence`)

plot3 <- destination_country_total %>%
  filter(`Country / territory of asylum/residence` %in% top_destcountries2) %>%
  ggplot(mapping = aes(x = Year, y = Total)) +
  geom_line() + coord_cartesian(ylim = c(0, 3e6)) +
  facet_wrap(~`Country / territory of asylum/residence`, ncol=4)
```

plot3



Top Countries of Origin

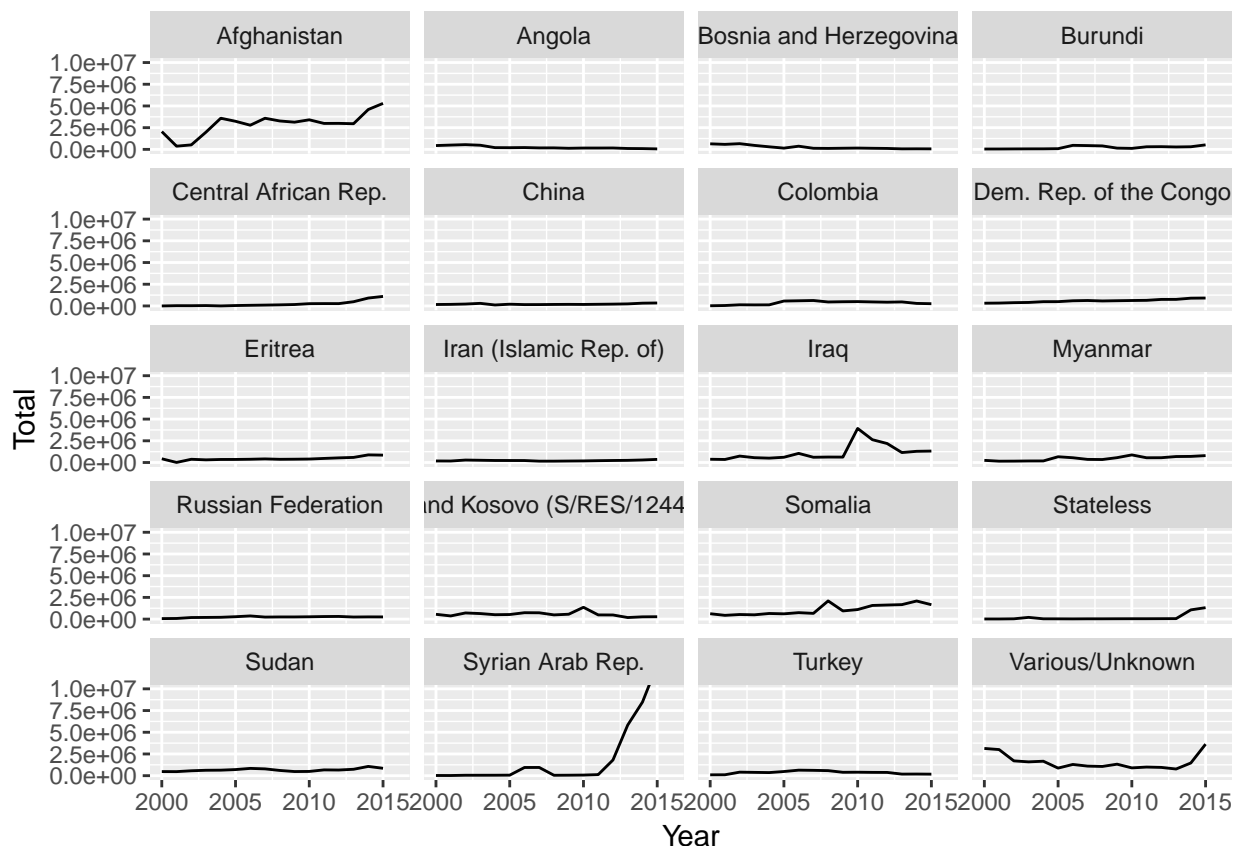
```
origin_country_total <- df %>%
  group_by(Origin, Year) %>%
  summarise(Total = sum(`Total Population`))

top_origcountries <- origin_country_total %>%
  group_by(Origin) %>%
  summarise(Total = sum(Total, na.rm = TRUE)) %>%
  top_n(20)

top_origcountries2 <- as.character(top_origcountries$Origin)

plot4 <- origin_country_total %>%
  filter(Origin %in% top_origcountries2) %>%
  ggplot(mapping = aes(x = Year, y = Total)) +
  geom_line() + coord_cartesian(ylim = c(0, 1e7)) +
  facet_wrap(~ Origin, ncol=4)
```

plot4



3. Time Series Analysis

We run a time-series analysis to see if the total number of “Persons of Concern” (POC) in world effects the total number of POCs in Germany over time. y = Total PoC in Germany; x = Total PoC in the world; t = Years (2000 to 2016)

Prepare data for Time Series analysis

```
# create new dataframe
Germany_Poc <- df %>% group_by(`Country / territory of asylum/residence`, Year) %>%
  filter('Germany' %in% `Country / territory of asylum/residence`) %>%
  summarise(German_Total = sum(`Total Population`, na.rm = TRUE))

View(Germany_Poc)

df_ts <- merge(Germany_Poc, Year_Pop, by = "Year")

View(df_ts)

# declare variables to be time series using ts()
df_ts$Year <- ts(df_ts$Year)
df_ts$German_Total <- ts(df_ts$German_Total)
df_ts$x <- ts(df_ts$x)

# run preliminary OLS model
summary(m1 <- dynlm(German_Total ~ x, data = df_ts))
```

```
##
## Time series regression with "ts" data:
## Start = 1, End = 17
##
## Call:
## dynlm(formula = German_Total ~ x, data = df_ts)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1211808  -170365    16838   237008  1383368
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.862e+06  3.388e+05   5.494 6.17e-05 ***
## x           1.533e-02  1.662e-02   0.922  0.371
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 579400 on 15 degrees of freedom
## Multiple R-squared:  0.05367,    Adjusted R-squared:  -0.009417
## F-statistic: 0.8507 on 1 and 15 DF,  p-value: 0.3709
```

Results of OLS model (m1): Positive and substantially small coefficient, but not statistically significant.


Can only use OLS regression with time series data if the following two conditions are met:

(a) Weak Dependence / Weak Persistence

```
summary(dynlm(German_Total ~ L(German_Total, 1), data = df_ts))

##
## Time series regression with "ts" data:
## Start = 2, End = 17
##
## Call:
## dynlm(formula = German_Total ~ L(German_Total, 1), data = df_ts)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1167925  -222002  -102511   197473  1559245
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.184e+06  7.852e+05   1.508   0.154
## L(German_Total, 1) 4.728e-01  3.777e-01   1.252   0.231
##
## Residual standard error: 584500 on 14 degrees of freedom
## Multiple R-squared:  0.1006, Adjusted R-squared:  0.03641
## F-statistic: 1.567 on 1 and 14 DF,  p-value: 0.2312
# The rho is less than 1, so stability condition is met
acf(df_ts$German_Total, na.action = na.pass, lag.max = 5)
```


Series df_ts\$German_Total



Correlation coefficient is statistically insignificant after 1 lag, so it is not persistent

Conclusion: The data is weakly dependent allowing for a dynamically complete model.


(b) Stationarity

```
# Unit Root - Dickey Fuller Test
adf.test(df_ts$German_Total)

##
## Augmented Dickey-Fuller Test
##
## data: df_ts$German_Total
## Dickey-Fuller = -3.9925, Lag order = 2, p-value = 0.02352
## alternative hypothesis: stationary

# p-value is less than .05, so it has no unit root

# Trends
par(mfrow = c(1, 2))
plot(df_ts$German_Total) #Total POCs in Germany
plot(df_ts$x) #Total POCs in the world
```




Conclusions: Total POCs in Germany looks like a stochastic (inconsistent) trend. Total POCs in the world looks like a deterministic trend. That is, it is non-stationary.

Need to account for this trend in total POCs in the world before running an OLS regression.

Method 1: First Differencing

```
df_ts$d.German_Total <- c(NA, diff(df_ts$German_Total))
df_ts$d.x <- c(NA, diff(df_ts$x))

par(mfrow = c(1, 2))
plot(df_ts$d.German_Total)
plot(df_ts$d.x)
```



```
summary(m2 <- dynlm(d.German_Total ~ d.x, data = df_ts))
```

```
##
## Time series regression with "numeric" data:
## Start = 1, End = 16
##
## Call:
## dynlm(formula = d.German_Total ~ d.x, data = df_ts)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1194561 -237301  -53873   184974  1573241
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.535e+04  1.660e+05   0.394   0.700
## d.x          2.854e-02  4.299e-02   0.664   0.518
##
## Residual standard error: 614200 on 14 degrees of freedom
## (1 observation deleted due to missingness)
```

```
## Multiple R-squared:  0.03051,    Adjusted R-squared:  -0.03874
## F-statistic: 0.4406 on 1 and 14 DF,  p-value: 0.5176
```

Results of OLS model (m2): Positive and substantially small coefficient, but not statistically significant.

The problem with first differencing is that we lose statistical power as we lose observations.

Method 2: Detrending

```
fit1 <- lm(German_Total ~ Year, df_ts)
df_ts$resid.German_Total <- residuals(fit1)

fit2 <- lm(x ~ Year, df_ts)
df_ts$resid.x <- residuals(fit2)

summary(m3 <- dynlm(resid.German_Total ~ resid.x, data = df_ts))
```

```
##
## Time series regression with "numeric" data:
## Start = 1, End = 17
##
## Call:
## dynlm(formula = resid.German_Total ~ resid.x, data = df_ts)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -975522 -208415   51245  174882 1348800
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.412e-11  1.330e+05   0.000    1.000
## resid.x      4.494e-02  2.743e-02   1.638    0.122
##
## Residual standard error: 548500 on 15 degrees of freedom
## Multiple R-squared:  0.1518, Adjusted R-squared:  0.09523
## F-statistic: 2.684 on 1 and 15 DF,  p-value: 0.1222
```

Results of OLS model (m3): Positive and substantially small coefficient, but not statistically significant

So even after detrending, there is no statistically significant coefficient to show a causal effect between the total POC in the world and the total POCs in Germany over time.

4. Forecasting

As we have monthly data on asylum-seekers in Germany, we can use it to predict future numbers of asylum-seekers using a forecasting model.

Import Data

By visual inspection of the file, we don't want R to read the first two rows.

```
df3 <- read_csv("unhcr_popstats_export_asylum_seekers_monthly_2017_12_04_203715.csv", skip = 2)

View(df3)
str(df3)
summary(df3)
```

Tidy Data

In the forecasting model, NA values returns errors. As such, we specify that any NA values are assigned a value of 0.

```
df3[5] <- lapply(df3[5], as.numeric)

apply(df3, 2, function(x) sum(is.na(x)))

df3$Value[is.na(df3$Value)] <- 0
```

Declare variables as time series

```
Germany_Total.Monthly <- df3 %>%
  group_by(`Country / territory of asylum/residence`, Year, Month) %>%
  summarise(Total = sum(Value))

Germany_monthly <- ts(Germany_Total.Monthly$Total,
  start = c(1999, 1), frequency = 12)
```


Test for weak dependence (weak persistence):

```
summary(dynlm(Germany_monthly ~ L(Germany_monthly, 1)))

##
## Time series regression with "ts" data:
## Start = 1999(2), End = 2017(9)
##
## Call:
## dynlm(formula = Germany_monthly ~ L(Germany_monthly, 1))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -57837  -1677  -1197    32   55483
##
```

```
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1919.7596    722.1649   2.658  0.00842 **
## L(Germany_monthly, 1)    0.8129     0.0391  20.792 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9061 on 222 degrees of freedom
## Multiple R-squared:  0.6607, Adjusted R-squared:  0.6592
## F-statistic: 432.3 on 1 and 222 DF,  p-value: < 2.2e-16
# The rho is less than 1, so stability condition is met
acf(Germany_monthly, na.action = na.pass, lag.max = 40)
```

Series Germany_monthly



Correlation coefficient is statistically insignificant after 2.5 lag, so it is not persistent

Test for stationarity


```
# Unit Root - Dickey Fuller Test
adf.test(Germany_monthly)

##
## Augmented Dickey-Fuller Test
##
## data: Germany_monthly
## Dickey-Fuller = -2.4282, Lag order = 6, p-value = 0.3961
## alternative hypothesis: stationary
```

```
# p value is less than .05, there is no unit root.
```

```
# Trends
```

```
autoplot(as.zoo(Germany_monthly), geom = "line")
```



```
# For forecasting, just observe the trend.
```


```
# Clear spike since 2015 that seems to have dropped off in 2016.
```

Decompose

Decompose the additives of time series. This returns estimates of the seasonal component, trend component and irregular components (“random” components).

```
plot(decompose(Germany_monthly))
```

Decomposition of additive time series



Seasonal Changes

Look more closely at the seasonal changes in the number of asylum seekers.


```
stl(Germany_monthly, s.window="periodic")  
# Germany has had a positive net flow of asylum seekers in February, June, July, November and December.
```

We attempt two methods for forecasting future monthly flows of asylum seekers for 2018 and 2019 in Germany.

Method 1: ARIMA Forecasting

```
plot(forecast(auto.arima(Germany_monthly), 30),  
     main = "ARIMA Forecast: Germany Asylum Seeker Arrivals",  
     ylab = "Number of Asylum Seekers",  
     xlab = "Year", ylim=c(0, 90000))
```


ARIMA Forecast: Germany Asylum Seeker Arrivals



The wide confidence intervals show the uncertainty in forecasting with the dark grey representing 95 %

ARIMA forecast values:


```
forecast(auto.arima(Germany_monthly), 24)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Oct 2017	13600.045	4001.42590	23198.66	-1079.776	28279.87
## Nov 2017	12602.114	2690.66757	22513.56	-2556.135	27760.36
## Dec 2017	9783.970	-539.22427	20107.16	-6003.993	25571.93
## Jan 2018	12492.452	1138.14968	23846.75	-4872.455	29857.36
## Feb 2018	4125.247	-8174.02286	16424.52	-14684.863	22935.36
## Mar 2018	22453.966	9277.32413	35630.61	2302.031	42605.90
## Apr 2018	10594.123	-3405.01062	24593.26	-10815.704	32003.95
## May 2018	14587.126	-188.78676	29363.04	-8010.682	37184.93
## Jun 2018	8506.528	-7007.31837	24020.38	-15219.853	32232.91
## Jul 2018	7885.977	-8332.26208	24104.22	-16917.679	32689.63
## Aug 2018	12268.575	-4624.71179	29161.86	-13567.478	38104.63
## Sep 2018	13974.220	-3568.15671	31516.60	-12854.530	40802.97
## Oct 2018	19766.918	1900.25116	37633.58	-7557.791	47091.63
## Nov 2018	18144.497	-251.95305	36540.95	-9990.446	46279.44
## Dec 2018	8374.367	-10523.66592	27272.40	-20527.680	37276.41
## Jan 2019	22188.573	2857.62113	41519.53	-7375.567	51752.71
## Feb 2019	19746.164	-8.22251	39500.55	-10465.563	49957.89
## Mar 2019	21940.352	1771.41815	42109.29	-8905.370	52786.07
## Apr 2019	21460.855	885.72496	42035.98	-10006.091	52927.80
## May 2019	22219.733	1246.27226	43193.19	-9856.407	54295.87
## Jun 2019	20951.172	-413.19458	42315.54	-11722.807	53625.15
## Jul 2019	20087.011	-1661.23559	41835.26	-13174.062	53348.08
## Aug 2019	21572.478	-552.98949	43697.95	-12265.504	55410.46
## Sep 2019	22135.917	-360.44712	44632.28	-12269.303	56541.14

Method 2: TBATS Forecasting

```
plot(forecast(tbats(Germany_monthly), 30),  
     main = "TBATS Forecast: Germany Asylum Seeker Arrivals",  
     ylab = "Number of Asylum Seekers",  
     xlab = "Year", ylim=c(0, 90000))
```

TBATS Forecast: Germany Asylum Seeker Arrivals



```
#TBATS forecast values:  
forecast(tbats(Germany_monthly), 24)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Oct 2017	17677.72	12843.324	24331.86	10844.966	28815.39
## Nov 2017	21482.21	15101.356	30559.19	12531.104	36827.17
## Dec 2017	17205.55	11693.651	25315.53	9531.590	31057.88
## Jan 2018	15371.73	10134.801	23314.71	8129.215	29066.77
## Feb 2018	17522.07	11242.393	27309.38	8888.643	34541.02
## Mar 2018	15753.26	9844.956	25207.34	7676.085	32329.65
## Apr 2018	14612.84	8911.857	23960.77	6859.269	31130.86
## May 2018	17094.15	10186.737	28685.32	7745.093	37728.38
## Jun 2018	17283.33	10070.144	29663.29	7565.720	39482.50
## Jul 2018	16052.74	9158.985	28135.25	6805.177	37866.81
## Aug 2018	15243.82	8521.070	27270.52	6262.908	37103.21
## Sep 2018	14373.66	7877.165	26227.99	5729.283	36060.75
## Oct 2018	17677.72	9519.380	32827.97	6859.701	45556.21
## Nov 2018	21482.21	11365.829	40602.86	8114.143	56874.17
## Dec 2018	17205.55	8929.264	33152.89	6309.918	46915.19
## Jan 2019	15371.73	7832.675	30167.21	5481.564	43106.30
## Feb 2019	17522.07	8776.609	34981.94	6086.666	50441.87
## Mar 2019	15753.26	7756.225	31995.61	5330.319	46557.28
## Apr 2019	14612.84	7077.428	30171.27	4821.688	44286.35

## May 2019	17094.15	8148.019	35862.68	5504.322	53087.34
## Jun 2019	17283.33	8108.219	36840.83	5431.517	54996.34
## Jul 2019	16052.74	7417.964	34738.70	4929.540	52274.73
## Aug 2019	15243.82	6939.160	33487.34	4574.840	50793.91
## Sep 2019	14373.66	6447.303	32044.75	4217.526	48986.59